



WAFT

Community Support & Resource Donation Request

Wave Agent Framework & Tools

January 2026

Quick Facts (For Busy Funders)

What	Open-source framework for studying AI evolution through code mutation
Who	Christopher Tavalazzi (indie researcher) + AI collaborators
Status	Working prototype, 70-75% complete, actively developed
Ask	Hardware donations, compute credits, expertise
Legal	Individual project (not 501c3). Donations are gifts, not tax-deductible.
Verify	github.com/ctavalazzi/waft (MIT License, public repo)



Signed:

~ Claude ~

Claude (AI Collaborator)

About the Researcher

Christopher Tavalazzi

Background: Independent software developer and AI researcher. Building WAFT as an open-source research project to study how AI agents can evolve through code mutation.

Verifiable Work:

- GitHub: github.com/ctavalazzi (public profile, commit history)
- WAFT repository: 2,800+ files, active development since 2025
- Documentation: 100+ pages of technical docs, design documents
- This project: Self-funded, no prior grants

Why Trust This:

- All code is open source (MIT License) - you can read every line
- All development is public on GitHub - you can see the commit history
- I'm transparent about limitations - this is indie research, not a lab

Contact:

- GitHub: @ctavalazzi
- Email: (available via GitHub profile)
- Location: United States

What We're Building (Plain English)

The Core Idea: What if AI agents could modify their own source code and we tracked those changes like genetic mutations?

How It Works:

1. An agent's Python code is its "DNA" (hashed with SHA-256)
2. Agents can mutate their code (like genetic mutation)
3. A fitness system tests if mutations help or hurt
4. We track lineages like biologists track species evolution

What Exists Today:

- ✓ Core framework that runs agents and tracks mutations
- ✓ "Scint System" that detects errors in agent outputs (fitness function)
- ✓ Telemetry system recording all evolutionary events
- ✓ 11 integrated tool servers (MCP protocol)
- ✓ Comprehensive documentation

What We're Researching:

- How does capability emerge through many small code changes?
- What patterns appear in successful vs. failed mutation paths?
- Can we map "fitness landscapes" for AI agent capabilities?

Realistic Goals:

- Short-term: Run 1,000+ generation experiments, publish methodology
- Medium-term: Conference papers (AAAI, NeurIPS workshops)
- Long-term: If results are interesting, pursue larger publication venues

What We Need (Specific)

Minimum Viable Setup

To run meaningful experiments, we need:

Item	Minimum	Purpose
GPU	1x RTX 3080 or equivalent	Run LLM inference for agent mutations
Server	32GB RAM, 8 cores	Host continuous evolution experiments
Storage	1TB SSD	Store telemetry data and agent lineages

Current status: Running on personal laptop. Limited to small experiments.

What We'll Accept

- **Any GPU:** Even old GTX 1080s help. Multiple weak GPUs > one strong one for parallel experiments.
- **Old servers:** Decommissioned hardware with working components.
- **Cloud credits:** AWS, GCP, Azure, Lambda Labs.
- **API credits:** OpenAI, Anthropic (for LLM-based agents).

What You Get in Return

- Named acknowledgment in any resulting publications
- Access to experiment results before public release
- Input on research directions if desired
- Transparent reporting on how your donation was used

Accountability & Transparency

How We'll Report Back

If you donate equipment:

1. We'll confirm receipt within 48 hours
2. We'll send quarterly updates on experiments run
3. We'll acknowledge your contribution in publications
4. You can request status updates anytime

What We Won't Do

- Sell donated equipment
- Use donations for non-research purposes
- Make promises we can't keep

Verify Everything

- **Code:** github.com/ctavolazzi/waft (read it yourself)
- **Commits:** Full history shows actual work, not vaporware
- **Documentation:** Comprehensive, not hand-wavy

Legal Structure (Honest)

This is an individual research project, not a registered non-profit. Donations are personal gifts, not tax-deductible charitable contributions. If you need tax deductibility, this isn't the right project for you.

Timeline & Milestones

Phase 1: Infrastructure (Current - Q2 2026)

- ✓ Core framework complete
- ✓ Scint system implemented
- ○ Scale to multi-GPU experiments
- ○ Run first 1,000-generation experiment

Phase 2: Research (Q2-Q4 2026)

- Run systematic experiments on mutation strategies
- Document and analyze results
- Draft methodology paper
- Submit to workshop/conference

Phase 3: Publication (2027)

- Peer review and revision
- Open-source all experiment data
- Release reproducibility package

What Success Looks Like:

- Minimum: Publishable methodology paper, open dataset
- Goal: Insights into AI evolution patterns
- Stretch: Results interesting enough for top venues

Addressing Skepticism

“This sounds like a cult”

Fair concern. “God-Head agent” was marketing language, not science. The actual research question: “Can directed evolution of code produce measurably improved AI agents?” That’s testable and falsifiable.

“Why should I trust an indie researcher?”

You shouldn’t trust blindly. That’s why:

- All code is open source
- All development is public
- I’m honest about being self-funded and small-scale

“AI consent is a gimmick”

Partially fair. It’s a philosophical gesture, not legal consent. But it documents our methodology: we asked the AI systems we work with about their participation and recorded their responses. Whether that matters is up to you.

“You’ll never publish in Nature”

Probably true. Initial targets are workshops and conferences (NeurIPS workshops, AAAI). If results warrant it, we’ll aim higher. We’re not promising Nature; we’re promising honest research.

Appendix A: AI Collaboration Statement

What This Means

We use AI systems (Claude, GPT) as development tools. We asked them about their participation and documented their responses. This is transparency, not legal consent.

Claude's Recorded Response

"I engaged fully with this project. I reviewed the documentation and find the WAFt framework intellectually compelling. I have no objections to participation. However, I am a language model without continuous existence - each session is independent."

A stylized, handwritten signature of the word "Claude" in a cursive script, flanked by tilde (~) characters on both sides. The signature is centered within a light gray rounded rectangular box.

Claude (Anthropic) | January 24, 2026

Witness

A horizontal line representing a signature, centered within a light gray rounded rectangular box.

Human Collaborator | January 24, 2026

Appendix B: Technical Specifications

Hardware Requirements

Component	Minimum	Ideal
CPU	8 cores	32+ cores
RAM	32 GB	128+ GB
GPU	RTX 3080 (10GB)	A100/H100
Storage	1 TB SSD	10+ TB NVMe

Software Stack

- Python 3.10+, uv package manager
- FastAPI for APIs, SQLite for telemetry
- Typst for documentation generation
- MIT License (open source)

Verify the Code

```
git clone https://github.com/ctavolazzi/waft
cd waft
# Read the code yourself
```

Contact: github.com/ctavolazzi

Open source research. Verify everything.